

Workshop Prerequisites

- Windows, MacOS:
 - Python + Numpy + OpenCV 3.4.1 from pip:
 - Online: `$ pip install opencv_python==3.4.1.15`
 - Offline: follow Readme.md for you platform on flash drive
- Linux (Ubuntu 16 recommended):
 - `$ apt get install python python-numpy`
 - goo.gl/Y4cEdD



xperience / ai

OpenCV Everywhere


Anna Petrovicheva

Agenda

- Intro
- OpenCV everywhere
- OpenCV C++ && bindings
- OpenCV Python hands on
- IOT use case Java/Python + JavaScript all together
- Optimization

Xperience.ai

- Deep Learning for Computer Vision
- Optimized for Mobile and FPGA
 - Quantization and Pruning
- Full pipeline of model development
 - Getting the data right
 - Building the model
 - Porting to target hardware




Tutorial materials

Github repo: github.com/xperience-ai/oreilly_opencv_workshop2018


- ▷ [cxx](#)
- ▷ [data](#)
- ▷ [java](#)
- ▷ [js](#)
- ▷ [python-template](#)
- ▷ [python](#)
- ▷ [slides](#)



OpenCV Everywhere



JavaScript





OpenCV for Python

- Python interface for OpenCV is binding to native C++ code
- Generated from C++ interface automatically, supported by OpenCV team
- Available with Python pip (built with limited features support)
- Numpy is base library for containers and math primitives
- Python 2.x and Python 3.x are supported
- OpenCV provides two independent libraries for Python 2.x and Python 3.x

OpenCV: C++ & Python

C++	Python
using namespace cv	import cv2 # For both Python2 and Python3
cv::Mat, cv::Mat::dot	Numpy.array, numpy.array.dot
#include <imgproc/imgproc.hpp> cv::blur(src, dst)	cv2.blur(src, dst)
capture >> frame;	has_frame, frame = capture.read()

Python & C++ Documentation Together

Function Documentation

\$ blobFromImage() [1/2]

```
Mat cv::dnn::blobFromImage ( InputArray   image,
                             double       scalefactor = 1.0,
                             const Size &  size = Size(),
                             const Scalar & mean = Scalar(),
                             bool         swapRB = true,
                             bool         crop = true
                             )
```

Python:

```
retval = cv.dnn.blobFromImage( image[, scalefactor[, size[, mean[, swapRB[, crop]]]]])
```

Creates 4-dimensional blob from image. Optionally resizes and crops `image` from center, subtract `mean` values, scales values by `scalefactor`, swap Blue and Red channels.

Parameters

- image** input image (with 1-, 3- or 4-channels).
- size** spatial size for output image
- mean** scalar with mean values which are subtracted from channels. Values are intended to be in (mean-R, mean-G, mean-B) order if

<https://docs.opencv.org/3.4.1/>

OpenCV for Java



- Java interface for OpenCV is binding to native C++ code
- Generated from C++ interface automatically with manually implemented data container adapters, supported by OpenCV team
- Available for desktop platforms: Windows, Linux, Mac OS, etc
- Available for Android with Android SDK & NDK integration

OpenCV: C++ -> Java

C++	Java
<code>cv::Mat, cv::Mat::dot</code>	<code>org.opencv.core.Mat, Mat.dot</code>
<code>#include <imgproc/imgproc.hpp></code> <code>cv::blur(src, dst)</code>	<code>import org.opencv.imgproc.*;</code> <code>Imgproc.blur(src, dst)</code>
<code>capture >> frame;</code>	<code>capture.read(frame);</code> <code>// Or use platform specific like android.Camera</code>

Javadoc



The screenshot shows the top of the OpenCV website. The OpenCV logo is on the left, followed by the text "OpenCV" and "Open Source Computer Vision". A version dropdown menu shows "3.4.1". Below this is a navigation bar with links: "Main Page", "Related Pages", "Modules", "Namespaces", "Classes", "Files", "Examples", and "Java documentation". The "Java documentation" link is highlighted with a red rectangle. Below the navigation bar is a section titled "OpenCV modules".

All Classes

Packages

org.opencv.aruco
org.opencv.bgsegm
org.opencv.bioinspired
org.opencv.calib3d
org.opencv.core

All Classes

AdaptiveManifoldFilter
AgastFeatureDetector
AKAZE
Algorithm
AlignExposures
AlignMTB
ANN_MLP
ANN_MLP_ANNEAL
Aruco
AverageHash
BackgroundSubtractor
BackgroundSubtractorCNT
BackgroundSubtractorGMG
BackgroundSubtractorGSOC
BackgroundSubtractorKNN
BackgroundSubtractorLSBP
BackgroundSubtractorLSBPDesc
BackgroundSubtractorMOG
BackgroundSubtractorMOG2

OVERVIEW PACKAGE CLASS TREE DEPRECATED INDEX HELP

PREV NEXT FRAMES NO FRAMES

OpenCV Java documentation (3.4.1)

Packages

Package	Description
org.opencv.aruco	
org.opencv.bgsegm	
org.opencv.bioinspired	
org.opencv.calib3d	
org.opencv.core	
org.opencv.dnn	
org.opencv.face	
org.opencv.features2d	
org.opencv.highgui	
org.opencv.img_hash	
org.opencv.imgcodecs	

- Introduction
- OpenCV Tutorials

<https://docs.opencv.org/3.4.1/javadoc/index.html>




OpenCV JS

- OpenCV JS is transpiled from C++ code and has the same structure and functions naming
- Works as JS or WebAssembly library in browser, has integrations with <video> and Canvas
- No reference manual yet

C++	JavaScript
<code>cv::Mat, cv::Mat::dot</code>	<code>cv.Mat, cv.Mat.dot</code>
<code>#include <imgproc/imgproc.hpp></code> <code>cv::blur(src, dst)</code>	<code>utils.loadOpenCv(() => {</code> <code>cv.blur(src, dst)</code> <code>});</code>
<code>capture >> frame;</code>	<code>capture.read()</code>

Golang, Haskell, D, other

- External bindings developed by community
- Manual implementation
- Not supported by core OpenCV team



Hands-on Scenario for Python

- Camera preview, live streaming
- Basic image processing and classic CV primitives
- Face Detection with DNN
- Face Recognition with DNN

Workshop Prerequisites

- Windows, MacOS:
 - Python + Numpy + OpenCV 3.4.1 from pip:
 - Online: `$ pip install opencv_python==3.4.1.15`
 - Offline: follow Readme.md for you platform on flash drive
- Linux (Ubuntu 16 recommended):
 - `$ apt get install python python-numpy`
 - goo.gl/Y4cEdD



Bootstrap

- For Windows and MacOS:
 - nothing
- For Linux:
 - `$ export PYTHONPATH=<opencv_distro>/lib/python<version>/dist-packages:$PYTHONPATH`
 - `$ export LD_LIBRARY_PATH =<opencv_distro>/lib:$LD_LIBRARY_PATH`
- Test:
 - `import numpy`
 - `import cv2`

Step 0: Video decoding & Camera Access

```
source = cv2.VideoCapture(0)
win_name = 'Camera Preview'
cv2.namedWindow(win_name, cv2.WINDOW_NORMAL)
```

```
while cv2.waitKey(1) != 27: # Escape
    has_frame, frame = source.read()
    if not has_frame:
        break
    cv2.imshow(win_name, frame)
```

```
source.release()
cv2.destroyWindow(win_name)
```



See *python-template/Camera/Camera.py* in repository and on stick

Step 1: Camera Filters

```
feature_params = dict( maxCorners = 50, qualityLevel = 0.3,  
                        minDistance = 7, blockSize = 7 )
```

```
...
```

```
if image_filter == PREVIEW:
```

```
    result = frame;
```

```
elif image_filter == CANNY:
```

```
    result = cv2.Canny(frame, 80, 90);
```

```
elif image_filter == BLUR:
```

```
    result = cv2.blur(frame, (13, 13));
```

```
elif image_filter == FEATURES:
```

```
    result = frame
```



```
    frame_gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
```

```
    corners = cv2.goodFeaturesToTrack(frame_gray, **feature_params)
```

```
    if corners is not None:
```

```
        for x, y in numpy.float32(corners).reshape(-1, 2):
```

```
            cv2.circle(result, (x,y), 10, (0, 255, 0), 1)
```



See [python-template/CameraFilters/CameraFilters.py](#) in repository and on stick

Step 2: Face Detection Initialization

Model parameters

in_width = 300

in_height = 300

mean = [104, 117, 123]

conf_threshold = 0.7

```
net = cv2.dnn.readNetFromCaffe("../data/deploy.prototxt",  
                                "../data/res10_300x300_ssd_iter_140000_fp16.caffemodel")
```


See python/Camera/Camera.py in repository and on stick

Step 2: Per-frame Face Detection Inference in Python

```
frame_height = frame.shape[0]
frame_width = frame.shape[1]
blob = cv2.dnn.blobFromImage(frame, 1.0, (in_width, in_height),
                              mean, False, False)

net.setInput(blob)
detections = net.forward()

for i in range(detections.shape[2]):
    confidence = detections[0, 0, i, 2]
    if confidence > conf_threshold:
        x_left_bottom = int(detections[0, 0, i, 3] * frame_width)
        y_left_bottom = int(detections[0, 0, i, 4] * frame_height)
        x_right_top = int(detections[0, 0, i, 5] * frame_width)
        y_right_top = int(detections[0, 0, i, 6] * frame_height)
```




Step 2: Per-frame Face Detection Inference in C++

```
Mat inputBlob = dnn::blobFromImage(frame, inScaleFactor, Size(inWidth, inHeight), meanVal, false, false);  
net.setInput(inputBlob, "data");  
Mat detection = net.forward("detection_out");
```

```
Mat detectionMat(detection.size[2], detection.size[3], CV_32F, detection.ptr<float>());  
for(int i = 0; i < detectionMat.rows; i++)  
{  
    float confidence = detectionMat.at<float>(i, 2);  
  
    if(confidence > confidenceThreshold)  
    {  
        float xLeftBottom = max(0, detectionMat.at<float>(i, 3) * frame.cols);  
        float yLeftBottom = max(0, detectionMat.at<float>(i, 4) * frame.rows);  
        float xRightTop = min(result.cols-1, detectionMat.at<float>(i, 5) * frame.cols);  
        float yRightTop = min(result.rows-1, detectionMat.at<float>(i, 6) * frame.rows);  
    }  
}
```

See `cxx/FaceDetection/main.cpp` in repository and on stick



Step 3: Face Recognition with OpenCV DNN

```
class DnnRecognizer:
```


```
    def __init__(self, model_path='../data/openface.nn4.small2.v1.t7',  
                  model_mean = [0, 0, 0], model_in_size = (96, 96), model_scale = 1.0 / 255, conf_threshold = 0.6):  
        self.known_faces = dict()  
        self.model = cv2.dnn.readNetFromTorch(model_path)  
        self.mean = model_mean  
        self.scale = model_scale
```

```
    def _face2vec(self, face):  
        blob = cv2.dnn.blobFromImage(face, self.scale, self.in_size, self.mean, False, False)  
        self.model.setInput(blob)  
        vec = self.model.forward()  
        return vec  
        self.in_size = model_in_size  
        self.confidence = conf_threshold
```

See `python-template/FaceDetectionAndRecognition/DnnRecognizer.py` in repository and on stick

Step 3: Face Recognition with OpenCV DNN 2

```
def introduce(self, image, name):  
    self.known_faces[name] = self._face2vec(image)  
  
def recognize(self, image):  
    vec = self._face2vec(image)  
    best_match_name = 'unknown'  
    best_match_score = self.confidence  
    # NOTE: Replace items() method to iteritems() if you use Python2  
    for name, descriptor in self.known_faces.items():  
        score = vec.dot(descriptor.T)  
        if (score > best_match_score):  
            best_match_score = score  
            best_match_name = name  
    return best_match_name
```



IOT Use case: Online Recognition

- **Server: HTTP REST API with POST requests on Java or Python:**
 - `/introduce?name=[person name]` Body - JPEG or PNG image
 - `/recognize` Body - jpeg image
 - *See `python/server/server.py` and `java/server/` in repository and on stick*
- **Client:**
 - WEB site with OpenCV JS
 - *See `js/cascade_face_detection_remote_recognition.html` in repository and on stick*
 - Android Application with OpenCV for Android

Step 1: Face detection with OpenCV JS

```
let video = document.getElementById('videoInput');  
let src = new cv.Mat(video.height, video.width, cv.CV_8UC4);  
let dst = new cv.Mat(video.height, video.width, cv.CV_8UC4);  
let gray = new cv.Mat();  
let cap = new cv.VideoCapture(video);  
let faces = new cv.RectVector();  
let classifier = new cv.CascadeClassifier();  
const FPS = 30;
```

Step 2: Face detection with OpenCV JS

```
let begin = Date.now();
cap.read(src);
src.copyTo(dst);
cv.cvtColor(dst, gray, cv.COLOR_RGBA2GRAY, 0);
classifier.detectMultiScale(gray, faces, 1.1, 3, 0);
for (let i = 0; i < faces.size(); ++i)
{
    let face = faces.get(i);
    let point1 = new cv.Point(face.x, face.y);
    let point2 = new cv.Point(face.x + face.width, face.y + face.height);
    cv.rectangle(dst, point1, point2, [255, 0, 0, 255]);
}
cv.imshow('canvasOutput', dst);
if(faces.size() == 1 && !introducing)
{
    face_roi = src.roi(faces.get(0))
    cv.imshow('introducingFormFace', face_roi);
}
```

Optimization

- OpenCV contains SSE, AVX, AVX2 optimizations for x86 and NEON for ARM CPUs
- OpenCV provides Transparent API accelerated with KHRONOS OpenCL
- OpenCV includes set of cudaXXX modules with NVIDIA CUDA accelerated algorithms
- Deep Learning inference is optimized for CPU and GPU, different backends could be used:
 - OpenCV native
 - OpenCL path
 - [Intel Inference Engine \(part of Intel OpenVINO\)](#)

Performance: DNN backends

```
backends = (cv2.dnn.DNN_BACKEND_DEFAULT, cv2.dnn.DNN_BACKEND_HALIDE,  
cv2.dnn.DNN_BACKEND_INFERENCE_ENGINE, cv2.dnn.DNN_BACKEND_OPENCV)
```

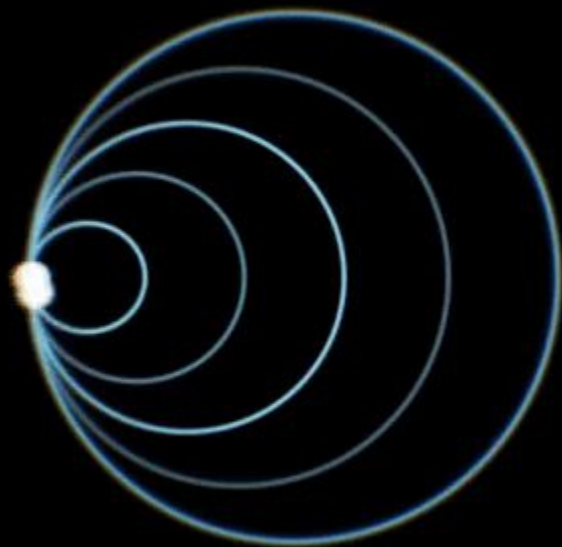
```
targets = (cv2.dnn.DNN_TARGET_CPU, cv2.dnn.DNN_TARGET_OPENCL,  
cv2.dnn.DNN_TARGET_OPENCL_FP16, cv2.dnn.DNN_TARGET_MYRIAD)
```

```
net.setPreferableBackend(backend)
```

```
net.setPreferableTarget(target)
```

Useful Links and References

- Workshop repo: https://github.com/xperience-ai/oreilly_opencv_workshop2018
- OpenCV @ Github: <https://github.com/opencv/opencv>
- OpenCV online documentation: <https://docs.opencv.org/3.4.1/>
- Q&A forum: <http://answers.opencv.org/questions/>



xperience.ai